

bank_fraud

Group Project

2023-11-09

1. Introduction

This project aims to identify fraudulent transactions with Credit Cards. Our objective is to build a Fraud detection system using Machine learning techniques. Each transaction is labelled either fraudulent or not fraudulent. Note that prevalence of fraudulent transactions is very low in the dataset, hence the dataset will be balanced using re-sampling techniques.

2. Data Understanding

The dataset contain 555719 transactions with 23 attributes that will be analyzed and then used to build a machine model for making predictions.

Dataset Breakdown: 23 attributes (22 predictive attributes and 1 goal field)

```
#importing Library

suppressMessages({
library(caret)
library(tidyverse)
library(randomForest)
library(corrplot)
library(rpart)
library(kernlab)
library(pROC)
library(ROSE) #Library for resampling
library(lubridate)
library(class)})

## Warning: package 'ROSE' was built under R version 4.3.2

#Load dataset
fraud_csv <- read.csv("fraud_dataset.csv")
head(fraud_csv)

##   X trans_date_trans_time      cc_num
merchant
## 1 0    2020-06-21 12:14:25 2.291164e+15      fraud_Kirlin and
Sons
## 2 1    2020-06-21 12:14:33 3.573030e+15      fraud_Sporer-
Keebler
## 3 2    2020-06-21 12:14:53 3.598215e+15 fraud_Swaniawski, Nitzsche and
```

```

Welch
## 4 3 2020-06-21 12:15:15 3.591920e+15 fraud_Haley
Group
## 5 4 2020-06-21 12:15:17 3.526826e+15 fraud_Johnston-
Casper
## 6 5 2020-06-21 12:15:37 3.040768e+13 fraud_Daugherty
LLC
## category amt first last gender
street
## 1 personal_care 2.86 Jeff Elliott M 351 Darlene
Green
## 2 personal_care 29.84 Joanne Williams F 3638 Marsh
Union
## 3 health_fitness 41.28 Ashley Lopez F 9333 Valentine
Point
## 4 misc_pos 60.05 Brian Williams M 32941 Krystal Mill Apt.
552
## 5 travel 3.19 Nathan Massey M 5783 Evan Roads Apt.
465
## 6 kids_pets 19.55 Danielle Evans F 76752 David Lodge Apt.
064
## city state zip lat long city_pop job
## 1 Columbia SC 29209 33.9659 -80.9355 333497 Mechanical engineer
## 2 Altonah UT 84002 40.3207 -110.4360 302 Sales professional, IT
## 3 Bellmore NY 11710 40.6729 -73.5365 34496 Librarian, public
## 4 Titusville FL 32780 28.5697 -80.8191 54767 Set designer
## 5 Falmouth MI 49632 44.2529 -85.0170 1126 Furniture designer
## 6 Breesport NY 14816 42.1939 -76.7361 520 Psychotherapist
## dob trans_num unix_time merch_lat
merch_long
## 1 1968-03-19 2da90c7d74bd46a0caf3777415b3ebd3 1371816865 33.98639 -
81.20071
## 2 1990-01-17 324cc204407e99f51b0d6ca0055005e7 1371816873 39.45050 -
109.96043
## 3 1970-10-21 c81755dbbba9d5c77f094348a7579be 1371816893 40.49581 -
74.19611
## 4 1987-07-25 2159175b9efe66dc301f149d3d5abf8c 1371816915 28.81240 -
80.88306
## 5 1955-07-06 57ff021bd3f328f8738bb535c302a31b 1371816917 44.95915 -
85.88473
## 6 1991-10-13 798db04aaceb4febd084f1a7c404da93 1371816937 41.74716 -
77.58420
## is_fraud
## 1 0
## 2 0
## 3 0
## 4 0
## 5 0
## 6 0

```

#checking initial structure and datatypes in dataset

```
str(fraud_csv)
```

```
## 'data.frame':    555719 obs. of  23 variables:
## $ X                : int  0 1 2 3 4 5 6 7 8 9 ...
## $ trans_date_trans_time: chr  "2020-06-21 12:14:25" "2020-06-21 12:14:33"
"2020-06-21 12:14:53" "2020-06-21 12:15:15" ...
## $ cc_num            : num  2.29e+15 3.57e+15 3.60e+15 3.59e+15
3.53e+15 ...
## $ merchant          : chr  "fraud_Kirlin and Sons" "fraud_Sporer-
Keebler" "fraud_Swaniawski, Nitzsche and Welch" "fraud_Haley Group" ...
## $ category          : chr  "personal_care" "personal_care"
"health_fitness" "misc_pos" ...
## $ amt               : num  2.86 29.84 41.28 60.05 3.19 ...
## $ first             : chr  "Jeff" "Joanne" "Ashley" "Brian" ...
## $ last              : chr  "Elliott" "Williams" "Lopez" "Williams" ...
## $ gender            : chr  "M" "F" "F" "M" ...
## $ street            : chr  "351 Darlene Green" "3638 Marsh Union"
"9333 Valentine Point" "32941 Krystal Mill Apt. 552" ...
## $ city              : chr  "Columbia" "Altonah" "Bellmore"
"Titusville" ...
## $ state             : chr  "SC" "UT" "NY" "FL" ...
## $ zip               : int  29209 84002 11710 32780 49632 14816 95528
57374 16858 76678 ...
## $ lat               : num  34 40.3 40.7 28.6 44.3 ...
## $ long              : num  -80.9 -110.4 -73.5 -80.8 -85 ...
## $ city_pop          : int  333497 302 34496 54767 1126 520 1139 343
3688 263 ...
## $ job               : chr  "Mechanical engineer" "Sales professional,
IT" "Librarian, public" "Set designer" ...
## $ dob              : chr  "1968-03-19" "1990-01-17" "1970-10-21"
"1987-07-25" ...
## $ trans_num         : chr  "2da90c7d74bd46a0caf3777415b3ebd3"
"324cc204407e99f51b0d6ca0055005e7" "c81755dbbba9d5c77f094348a7579be"
"2159175b9efe66dc301f149d3d5abf8c" ...
## $ unix_time         : int  1371816865 1371816873 1371816893 1371816915
1371816917 1371816937 1371816944 1371816950 1371816970 1371816971 ...
## $ merch_lat         : num  34 39.5 40.5 28.8 45 ...
## $ merch_long        : num  -81.2 -110 -74.2 -80.9 -85.9 ...
## $ is_fraud          : int  0 0 0 0 0 0 0 0 0 0 ...
```

#checking for cardinality of columns and also for null values

```
fraud_csv %>% summarise_all(n_distinct)
```

```
##      X trans_date_trans_time cc_num merchant category  amt first last
gender
## 1 555719                544760    924      693      14 37256   341  471
2
## street city state zip lat long city_pop job dob trans_num unix_time
merch_lat
```

```
## 1      924  849      50 912 910  910      835 478 910      555719      544760
546490
## merch_long is_fraud
## 1      551770      2

sum(is.na.data.frame(fraud_csv))

## [1] 0
```

Converting features to appropriate datatypes

```
#convert time to posix
fraud_csv$datetime <- as.POSIXct(fraud_csv$trans_date_trans_time, format="%Y-
%m-%d %H:%M:%S")

#extract dates
fraud_csv$date <- as.Date(fraud_csv$datetime)
#extract age from dob
fraud_csv$dob <- as.Date(fraud_csv$dob)
fraud_csv$age_2022 <- round(as.numeric(difftime(as.Date("2022-01-01"),
fraud_csv$dob, units = "days"))) / 365.25, 0)

#convert dates +times into day, month and hour
fraud_csv$time <- format(fraud_csv$datetime, format="%H:%M:%S")
fraud_csv$hour <- as.numeric(hour(fraud_csv$datetime))
fraud_csv$weekday <- weekdays(as.Date(fraud_csv$datetime))
fraud_csv$weekday2 <- as.numeric(format(fraud_csv$datetime, format = "%u"))
fraud_csv$month<- as.numeric(month(fraud_csv$datetime))

#convert gender to int
fraud_csv$gender <- ifelse(fraud_csv$gender=="M", 1, 0)

#concatenation of name
fraud_csv$full_name <- paste(fraud_csv$first, fraud_csv$last, sep = " ")

#checking new structure of dataset with new features created
str(fraud_csv)

## 'data.frame':    555719 obs. of  32 variables:
## $ X                : int  0 1 2 3 4 5 6 7 8 9 ...
## $ trans_date_trans_time: chr  "2020-06-21 12:14:25" "2020-06-21 12:14:33"
"2020-06-21 12:14:53" "2020-06-21 12:15:15" ...
## $ cc_num            : num  2.29e+15 3.57e+15 3.60e+15 3.59e+15
3.53e+15 ...
## $ merchant          : chr  "fraud_Kirlin and Sons" "fraud_Sporer-
Keebler" "fraud_Swaniawski, Nitzsche and Welch" "fraud_Haley Group" ...
## $ category          : chr  "personal_care" "personal_care"
"health_fitness" "misc_pos" ...
## $ amt               : num  2.86 29.84 41.28 60.05 3.19 ...
## $ first             : chr  "Jeff" "Joanne" "Ashley" "Brian" ...
## $ last              : chr  "Elliott" "Williams" "Lopez" "Williams" ...
```

```

## $ gender          : num  1 0 0 1 1 0 0 0 1 0 ...
## $ street          : chr   "351 Darlene Green" "3638 Marsh Union"
"9333 Valentine Point" "32941 Krystal Mill Apt. 552" ...
## $ city            : chr   "Columbia" "Altonah" "Bellmore"
"Titusville" ...
## $ state           : chr   "SC" "UT" "NY" "FL" ...
## $ zip             : int    29209 84002 11710 32780 49632 14816 95528
57374 16858 76678 ...
## $ lat             : num    34 40.3 40.7 28.6 44.3 ...
## $ long            : num   -80.9 -110.4 -73.5 -80.8 -85 ...
## $ city_pop        : int    333497 302 34496 54767 1126 520 1139 343
3688 263 ...
## $ job             : chr   "Mechanical engineer" "Sales professional,
IT" "Librarian, public" "Set designer" ...
## $ dob            : Date, format: "1968-03-19" "1990-01-17" ...
## $ trans_num       : chr   "2da90c7d74bd46a0caf3777415b3ebd3"
"324cc204407e99f51b0d6ca0055005e7" "c81755dbbba9d5c77f094348a7579be"
"2159175b9efe66dc301f149d3d5abf8c" ...
## $ unix_time       : int    1371816865 1371816873 1371816893 1371816915
1371816917 1371816937 1371816944 1371816950 1371816970 1371816971 ...
## $ merch_lat       : num    34 39.5 40.5 28.8 45 ...
## $ merch_long      : num   -81.2 -110 -74.2 -80.9 -85.9 ...
## $ is_fraud        : int    0 0 0 0 0 0 0 0 0 0 ...
## $ datetime        : POSIXct, format: "2020-06-21 12:14:25" "2020-06-
21 12:14:33" ...
## $ date            : Date, format: "2020-06-21" "2020-06-21" ...
## $ age_2022        : num    54 32 51 34 66 30 71 50 49 66 ...
## $ time            : chr   "12:14:25" "12:14:33" "12:14:53" "12:15:15"
...
## $ hour            : num    12 12 12 12 12 12 12 12 12 12 ...
## $ weekday         : chr   "Sunday" "Sunday" "Sunday" "Sunday" ...
## $ weekday2        : num    7 7 7 7 7 7 7 7 7 7 ...
## $ month           : num    6 6 6 6 6 6 6 6 6 6 ...
## $ full_name       : chr   "Jeff Elliott" "Joanne Williams" "Ashley
Lopez" "Brian Williams" ...

write.csv(fraud_csv, file = "fraud_csv_new.csv")

#counting number of occurrences of fraud(1) and no_fraud(0)
table(fraud_csv$is_fraud)

##
##          0          1
## 553574    2145


prop.table(table(fraud_csv$is_fraud))

##
##          0          1
## 0.996140136 0.003859864

```

```
# visualizing data imbalance
ggplot(fraud_csv, aes(x = is_fraud)) +
  geom_bar() +
  geom_text(stat='count', aes(label=..count..), vjust=-1)

## Warning: The dot-dot notation (`..count..`) was deprecated in ggplot2
## 3.4.0.
## i Please use `after_stat(count)` instead.
## This warning is displayed once every 8 hours.
## Call `lifecycle::last_lifecycle_warnings()` to see where this warning was
## generated.
```



Extrapolating cost of Fraud (incidents and monetary value) for the whole year

```
#days_in_dataset: 194

days_in_dataset <- 194

# Calculate expected number of fraud cases per year
fraud_per_day <- sum(fraud_csv$is_fraud) / days_in_dataset
fraud_per_year <- fraud_per_day * 365
```

```

fraud_per_year_rounded <- round(fraud_per_year, 0)
fraud_per_year_rounded

## [1] 4036

# Calculate average non-fraudulent amount per day
nf_amount_df <- subset(fraud_csv, is_fraud == 0)
nf_amount_per_day <- sum(nf_amount_df$amt) / days_in_dataset
nf_amount_per_day_rounded <- round(nf_amount_per_day, 2)
paste(nf_amount_per_day_rounded, "$")

## [1] "192935.97 $"

# Calculate average fraudulent amount per day
fraud_amount_df <- subset(fraud_csv, is_fraud == 1)
fraud_amount_per_day <- sum(fraud_amount_df$amt) / days_in_dataset
fraud_amount_per_day_rounded <- round(fraud_amount_per_day, 2)
paste(fraud_amount_per_day_rounded, "$")

## [1] "5841.88 $"

# Estimate total fraudulent amount per year
total_fraud_yearly <- round(fraud_amount_per_day * 365, 2)
paste(total_fraud_yearly, "$")

## [1] "2132286.12 $"

# Estimate total non-fraudulent amount per year
total_non_fraud_yearly <- round(nf_amount_per_day * 365, 2)
paste(total_non_fraud_yearly, "$")

## [1] "70421629.52 $"

```

3. Exploratory Data Analysis

Before running machine learning models, features will be explored to see if there is any trends to point to prevalence of fraudulent transactions.

```


#sub-setting dataset for just fraudulent transactions
fraud_txns <- fraud_csv[fraud_csv$is_fraud == 1,]
dim(fraud_txns)

## [1] 2145    32

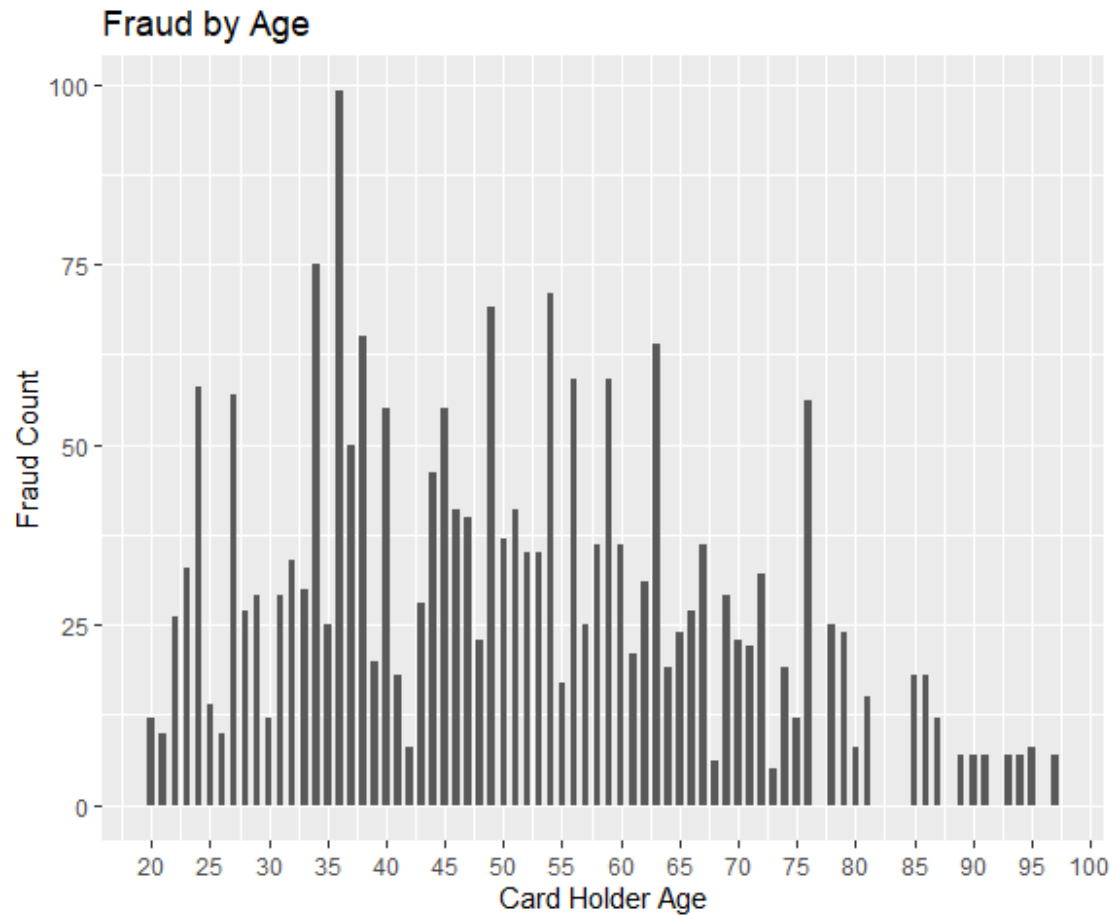
#checking distribution of fraudulent transaction amount
df_filtered <- fraud_txns[fraud_txns$amt < quantile(fraud_txns$amt, 0.99),]
hist(df_filtered$amt, breaks=100, main='Histogram of Fraud Transaction
Amounts', xlab='Amount')

```


Histogram of Fraud Transaction Amounts



```
fraud_by_age <- fraud_txns %>%  
  group_by(age_2022) %>%  
  summarise(count_fraud = n())  
  
ggplot(fraud_by_age, aes(age_2022, count_fraud, width =0.6)) +  
  geom_bar(stat = "identity") +  
  labs(title = "Fraud by Age", x = "Card Holder Age", y = "Fraud Count") +  
  scale_x_continuous(n.breaks=20)
```





```
fraud_by_hour <- fraud_txns %>%  
  group_by(hour) %>%  
  summarise(count_fraud = n())  
  
#visualizing fraud distribution by hour  
ggplot(fraud_by_hour, aes(hour, count_fraud)) +  
  geom_line() +  
  labs(title = "Fraud Transactions by the Hour",  
        x = "Hour",  
        y = "Number of Fraud Transactions")
```



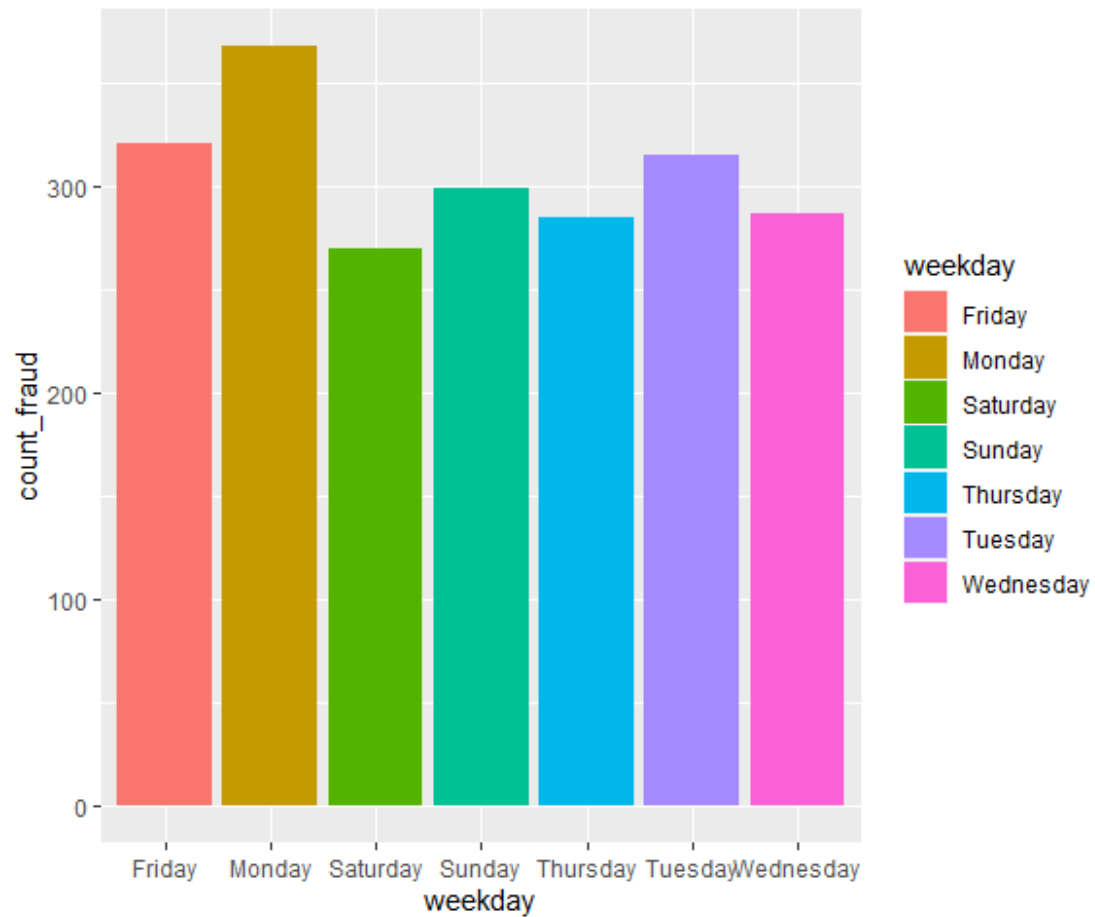
There is a clear trend of fraudulent transactions in the first 3 hours of the day and also last 2 hours.

```
fraud_by_category <- fraud_txns %>%  
  group_by(category) %>%  
  summarise(count_fraud = n())  
  
ggplot(data= fraud_by_category, aes(x=category, y= count_fraud))+  
  geom_bar(stat="identity")+  
  theme(axis.text.x = element_text(angle = 90, vjust = 0.5, hjust = 1))
```




Higher fraud rates in grocery_pos and shopping_net transactions

```
fraud_by_weekday <- fraud_txns %>%  
  group_by(weekday) %>%  
  summarise(count_fraud = n())  
  
ggplot(data= fraud_by_weekday, aes(x=weekday, y= count_fraud, fill=weekday))+  
  geom_bar(stat="identity")
```



```
boxplot(amt ~ is_fraud, data = fraud_csv,  
  main = "Fraud Status by Transaction Amount",  
  xlab = "Fraud Status",  
  ylab = "Transaction Amount",  
  names = c("No Fraud", "Fraud"))
```



4 Balancing Dataset and Feature Selection

Due to the imbalance that exists in the dataset, a new dataset will be created with a 60:40 balance for non-fraud(0) and fraud(1) cases. ML models will be trained on both the balanced and unbalanced dataset with testing done on the unbalanced (original) dataset.

4.1 Balancing Dataset

#setting the no of fraud as non_fraud cases as well as desired percentage of fraud in new dataset

```
r0 <- 0.40
```

```
n0 <- nrow(fraud_csv)
```

```
sampling_result <- ovun.sample(is_fraud ~ ., data = fraud_csv, method =  
"both", N = n0,
```

```
                                p = r0, seed = 2018)
```

```
balanced_data <- sampling_result$data
```

```
table(balanced_data$is_fraud)
```

```
##
```

```
##      0      1
```

```
## 333488 222231
```

```
prop.table(table(balanced_data$is_fraud))
```

```
##
```


```
##          0          1
```

```
## 0.6001019 0.3998981
```

```
ggplot(balanced_data, aes(x = is_fraud)) +
```

```
geom_bar() +
```

```
geom_text(stat='count', aes(label=..count..), vjust=-1)
```



4.3 Feature Selection

Feature Set 1 - All variables that do not contain relevant information and also modified variables will be dropped.

Balanced Data

Remove unnecessary features using subset

```
balanced_data2 <- subset(balanced_data, select = -c(X, trans_date_trans_time,
cc_num, merchant, first, last,
city, dob, trans_num, unix_time, datetime, date,
state, full_name))
```

```

str(balanced_data2)

## 'data.frame':    555719 obs. of  14 variables:
## $ category   : chr  "food_dining" "gas_transport" "personal_care"
## $ grocery_net : chr  "grocery_net" ...
## $ amt        : num  65.2 55.9 73.6 47.2 23.8 ...
## $ gender      : num  1 0 1 1 0 0 1 1 1 1 ...
## $ zip         : int  74633 97014 16362 54559 29438 31563 29209 95827 11964
## $ lat         : num  36.7 45.7 41.5 46.5 32.5 ...
## $ long        : num  -96.8 -121.9 -79.9 -90.4 -80.3 ...
## $ city_pop    : int  471 1288 1102 795 2408 1324 333497 757530 1858 471 ...
## $ merch_lat   : num  37.4 45.2 41.2 45.5 32.9 ...
## $ merch_long  : num  -96.9 -121.6 -79.1 -90.8 -81 ...
## $ is_fraud    : int  0 0 0 0 0 0 0 0 0 0 ...
## $ age_2022    : num  81 86 32 36 24 95 54 43 38 81 ...
## $ hour        : num  22 7 15 8 23 14 9 8 19 14 ...
## $ weekday2    : num  2 2 4 7 7 2 6 5 2 5 ...
## $ month       : num  10 7 11 11 11 6 11 11 12 7 ...

#encoding categorical variables
dmy <- dummyVars(" ~ .", data = balanced_data2, fullRank = T)
balanced_data3 <- data.frame(predict(dmy, newdata = balanced_data2))

#convert numerical data that are categories to factors
balanced_data3$gender <- factor(balanced_data3$gender)
balanced_data3$month <- factor(balanced_data3$month)
balanced_data3$weekday2 <- factor(balanced_data3$weekday2)

str(balanced_data3)

## 'data.frame':    555719 obs. of  26 variables:
## $ categoryfood_dining : num  1 0 0 0 0 0 0 0 0 0 ...
## $ categorygas_transport : num  0 1 0 0 0 0 0 0 0 0 ...
## $ categorygrocery_net : num  0 0 0 1 0 0 0 0 0 0 ...
## $ categorygrocery_pos : num  0 0 0 0 0 0 0 1 0 0 ...
## $ categoryhealth_fitness: num  0 0 0 0 0 0 1 0 0 0 ...
## $ categoryhome : num  0 0 0 0 0 0 0 0 0 0 ...
## $ categorykids_pets : num  0 0 0 0 0 0 0 0 0 0 ...
## $ categorymisc_net : num  0 0 0 0 0 0 0 0 0 0 ...
## $ categorymisc_pos : num  0 0 0 0 0 0 1 0 0 0 ...
## $ categorypersonal_care : num  0 0 1 0 0 0 0 0 0 0 ...
## $ categoryshopping_net : num  0 0 0 0 1 0 0 0 0 0 ...
## $ categoryshopping_pos : num  0 0 0 0 0 0 0 0 0 0 ...
## $ categorytravel : num  0 0 0 0 0 0 0 0 0 1 ...
## $ amt : num  65.2 55.9 73.6 47.2 23.8 ...
## $ gender : Factor w/ 2 levels "0","1": 2 1 2 2 1 1 2 2 2 2 ...
## $ zip : num  74633 97014 16362 54559 29438 ...

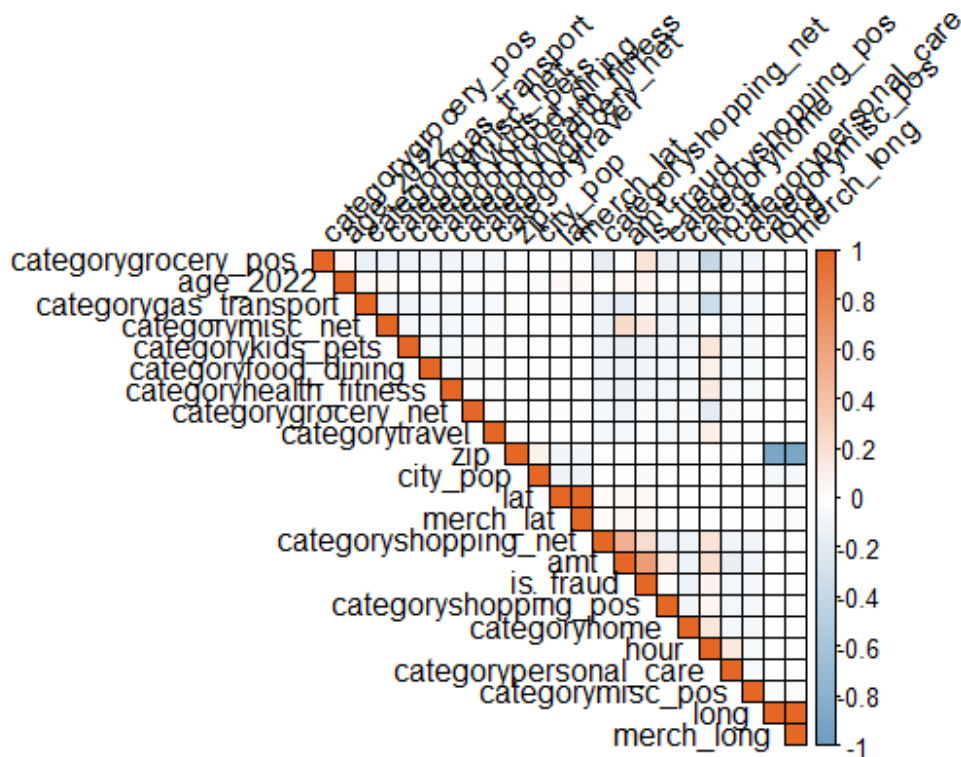
```

```
## $ lat : num 36.7 45.7 41.5 46.5 32.5 ...
## $ long : num -96.8 -121.9 -79.9 -90.4 -80.3 ...
## $ city_pop : num 471 1288 1102 795 2408 ...
## $ merch_lat : num 37.4 45.2 41.2 45.5 32.9 ...
## $ merch_long : num -96.9 -121.6 -79.1 -90.8 -81 ...
## $ is_fraud : num 0 0 0 0 0 0 0 0 0 ...
## $ age_2022 : num 81 86 32 36 24 95 54 43 38 81 ...
## $ hour : num 22 7 15 8 23 14 9 8 19 14 ...
## $ weekday2 : Factor w/ 7 levels "1","2","3","4",...: 2 2 4 7
7 2 6 5 2 5 ...
## $ month : Factor w/ 7 levels "6","7","8","9",...: 5 2 6 6
6 1 6 6 7 2 ...
```

#Checking for multicollinearity among variables using correlation matrix

```
cor_matrix <- cor(balanced_data3[, sapply(balanced_data3, is.numeric)])

corrplot(cor_matrix, method = "color", type = "upper", order = "hclust",
          tl.col = "black", tl.srt = 45,
          # Add grid and color
          addgrid.col = "black", col = colorRampPalette(c("#6D9EC1", "white",
"#E46726"))(200))
```

###

Unbalanced data

```
# Remove unnecessary features using subset
unbalanced_data <- subset(fraud_csv, select = -c(X, trans_date_trans_time,
cc_num,merchant, first, last,
                                street,
city,dob,trans_num,unix_time, datetime, date,
                                time, weekday, job,
state, full_name))

#encoding categorical variables
dmy <- dummyVars(" ~ .", data = unbalanced_data, fullRank = T)
unbalanced_data2 <- data.frame(predict(dmy, newdata = unbalanced_data))

#convert numerical data that are categories to factors
unbalanced_data2$gender <- factor(unbalanced_data2$gender)
unbalanced_data2$month <- factor(unbalanced_data2$month)
unbalanced_data2$weekday2 <- factor(unbalanced_data2$weekday2)

str(unbalanced_data2)
```

```
## 'data.frame': 555719 obs. of 26 variables:
## $ categoryfood_dining : num 0 0 0 0 0 0 0 0 0 1 ...
## $ categorygas_transport : num 0 0 0 0 0 0 0 0 0 0 ...
## $ categorygrocery_net : num 0 0 0 0 0 0 0 0 0 0 ...
## $ categorygrocery_pos : num 0 0 0 0 0 0 0 0 0 0 ...
## $ categoryhealth_fitness: num 0 0 1 0 0 0 1 0 0 0 ...
## $ categoryhome : num 0 0 0 0 0 0 0 0 0 0 ...
## $ categorykids_pets : num 0 0 0 0 0 1 0 0 0 0 ...
## $ categorymisc_net : num 0 0 0 0 0 0 0 0 0 0 ...
## $ categorymisc_pos : num 0 0 0 1 0 0 0 0 0 0 ...
## $ categorypersonal_care : num 1 1 0 0 0 0 0 1 0 0 ...
## $ categoryshopping_net : num 0 0 0 0 0 0 0 0 0 0 ...
## $ categoryshopping_pos : num 0 0 0 0 0 0 0 0 1 0 ...
## $ categorytravel : num 0 0 0 0 1 0 0 0 0 0 ...
## $ amt : num 2.86 29.84 41.28 60.05 3.19 ...
## $ gender : Factor w/ 2 levels "0","1": 2 1 1 2 2 1 1 1 2 1
...
## $ zip : num 29209 84002 11710 32780 49632 ...
## $ lat : num 34 40.3 40.7 28.6 44.3 ...
## $ long : num -80.9 -110.4 -73.5 -80.8 -85 ...
## $ city_pop : num 333497 302 34496 54767 1126 ...
## $ merch_lat : num 34 39.5 40.5 28.8 45 ...
## $ merch_long : num -81.2 -110 -74.2 -80.9 -85.9 ...
## $ is_fraud : num 0 0 0 0 0 0 0 0 0 0 ...
## $ age_2022 : num 54 32 51 34 66 30 71 50 49 66 ...
## $ hour : num 12 12 12 12 12 12 12 12 12 12 ...
## $ weekday2 : Factor w/ 7 levels "1","2","3","4",...: 7 7 7 7
7 7 7 7 7 7 ...
## $ month : Factor w/ 7 levels "6","7","8","9",...: 1 1 1 1
1 1 1 1 1 1 ...
```

5 Machine Learning Modelling

4 Models will be used to evaluate this classification task on both the balanced and unbalanced datasets. They are: i: Logistic Regression ii: Decision Tree iii: KNN iv: Random Forests

The models will be compared using classification ML metrics then the best one selected.

First the dataset will be split into Training and Test set

Setting set for Balanced dataset

#To achieve reproducible model; set the random seed number

```
set.seed(100)
```

Data is split into training and test set in a 80:20 ratio

```
TrainingIndex1 <- createDataPartition(balanced_data3$is_fraud, p=0.75, list = FALSE)
```

```

TrainingSet1 <- balanced_data3[TrainingIndex1,]# Training Set
TestingSet1 <- balanced_data3[-TrainingIndex1,]# Test Set

#To achieve reproducible model; set the random seed number
set.seed(1000)

# Data is split into training and test set in a 80:20 ratio
TrainingIndex2 <- createDataPartition(unbalanced_data2$is_fraud, p=0.75, list
= FALSE)

TrainingSet2 <- unbalanced_data2[TrainingIndex2,]# Training Set
TestingSet2 <- unbalanced_data2[-TrainingIndex2,]# Test Set

```

5.1 Logistic Regression

5.1.1 LR Balanced Dataset

```

cls <- glm(is_fraud~., family='binomial',data=TrainingSet1)

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

cut=0.5

#Calculate the training error
yhat = (predict(cls,TrainingSet1,type="response")>cut)
tr.err = mean(TrainingSet1$is_fraud != yhat)
tr.err

## [1] 0.1402217

#Calculate the testing error
yhat_test=(predict(cls,TestingSet2,type="response")>cut)
te_err=mean(TestingSet2$is_fraud!=yhat_test)
te_err

## [1] 0.05466101

# Prediction on TestingSet using Logistic Regression
cls_prediction <- predict(cls, TestingSet2, type ="response")
head(cls_prediction)

##          2          8          13          14          27          30
## 0.20767808 0.18632325 0.05975286 0.08778102 0.09441230 0.10792785

#Assigning probabilities - If prediction exceeds threshold of 0.5, 1 else 0
cls_prediction <- ifelse(cls_prediction >0.5,1,0)
head(cls_prediction)

##  2  8 13 14 27 30
##  0  0  0  0  0  0

```

```

#Computing confusion matrix values
confusionMatrix(factor(TestingSet2$is_fraud),factor(cls_prediction), mode
='everything', positive ="0")

## Confusion Matrix and Statistics
##
##              Reference
## Prediction      0      1
##      0 130926   7470
##      1    124    409
##
##              Accuracy : 0.9453
##              95% CI : (0.9441, 0.9465)
##      No Information Rate : 0.9433
##      P-Value [Acc > NIR] : 0.0004544
##
##              Kappa : 0.0907
##
##  Mcnemar's Test P-Value : < 2.2e-16
##
##              Sensitivity : 0.99905
##              Specificity : 0.05191
##              Pos Pred Value : 0.94602
##              Neg Pred Value : 0.76735
##              Precision : 0.94602
##              Recall : 0.99905
##              F1 : 0.97182
##              Prevalence : 0.94329
##              Detection Rate : 0.94240
##              Detection Prevalence : 0.99616
##              Balanced Accuracy : 0.52548
##
##              'Positive' Class : 0
##

```

5.1.2 Unbalanced Dataset

```

cls2 <- glm(is_fraud~., family='binomial',data=TrainingSet2)

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

cut=0.5

#Calculate the training error
yhat = (predict(cls2,TrainingSet2,type="response")>cut)
tr.err = mean(TrainingSet2$is_fraud != yhat)
tr.err

## [1] 0.004109983

#Calculate the testing error
yhat_test=(predict(cls2,TestingSet2,type="response")>cut)

```

```

te_err=mean(TestingSet2$is_fraud!=yhat_test)
te_err

## [1] 0.004081221

# Prediction on TestingSet using Logistic Regression
cls_prediction2 <- predict(cls2, TestingSet2, type ="response")
head(cls_prediction2)

##           2           8           13           14           27
30
## 0.0005146160 0.0007817741 0.0004053561 0.0018323282 0.0011362903
0.0006171092

#Assigning probabilities - If prediction exceeds threshold of 0.5, 1 else 0
cls_prediction2 <- ifelse(cls_prediction2 >0.5,1,0)
head(cls_prediction2)

##  2  8 13 14 27 30
##  0  0  0  0  0  0

#Computing confusion matrix values
confusionMatrix(factor(TestingSet2$is_fraud),factor(cls_prediction2), mode
='everything', positive ="0")

## Confusion Matrix and Statistics
##
##              Reference
## Prediction      0      1
##      0 138362      34
##      1   533       0
##
##              Accuracy : 0.9959
##              95% CI : (0.9956, 0.9962)
##      No Information Rate : 0.9998
##      P-Value [Acc > NIR] : 1
##
##              Kappa : -5e-04
##
##  McNemar's Test P-Value : <2e-16
##
##              Sensitivity : 0.9962
##              Specificity : 0.0000
##      Pos Pred Value : 0.9998
##      Neg Pred Value : 0.0000
##              Precision : 0.9998
##              Recall : 0.9962
##              F1 : 0.9980
##              Prevalence : 0.9998
##      Detection Rate : 0.9959
##      Detection Prevalence : 0.9962

```

```
##          Balanced Accuracy : 0.4981
##
##          'Positive' Class : 0
##
```

5.2 Decision Trees

5.2.1 DT Balanced Dataset

```
tree1 <- rpart(is_fraud ~., method = 'class', data = TrainingSet1, control =
rpart.control(cp = 0.0001))
```

```
# Predict using the decision tree model on the test data
```

```
predicted_values <- predict(tree1, newdata = TestingSet2, type = "class")
```

```
# Assuming 'is_fraud' is the actual target variable in the test data
```

```
actual_values <- TestingSet2$is_fraud
```

```
actual_values <- as.factor(actual_values)
```

```
levels(actual_values) <- c("0", "1")
```

```
tree_predictions1 <- predict(tree1, TestingSet2, type = 'class')
```

```
head(tree_predictions1)
```

```
##  2  8 13 14 27 30
```

```
##  0  0  0  0  0  0
```

```
## Levels: 0 1
```

```
# Create confusion matrix
```

```
conf_matrix1 <- confusionMatrix(predicted_values, actual_values)
```

```
print(conf_matrix1)
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

```
## Prediction      0      1
```

```
##           0 137125      0
```

```
##           1  1271    533
```

```
##
```

```
##           Accuracy : 0.9909
```

```
##           95% CI : (0.9903, 0.9913)
```

```
##       No Information Rate : 0.9962
```

```
##       P-Value [Acc > NIR] : 1
```

```
##
```

```
##           Kappa : 0.4529
```

```
##
```

```
## Mcnemar's Test P-Value : <2e-16
```

```
##
```

```
##           Sensitivity : 0.9908
```

```
##           Specificity : 1.0000
```

```
##       Pos Pred Value : 1.0000
```

```
##       Neg Pred Value : 0.2955
```

```
##           Prevalence : 0.9962
##           Detection Rate : 0.9870
##      Detection Prevalence : 0.9870
##           Balanced Accuracy : 0.9954
##
##           'Positive' Class : 0
##
```

5.2.2 DT Unbalanced Dataset

```
tree2 <- rpart(is_fraud ~., method = 'class', data = TrainingSet2, control =
rpart.control(cp = 0.0001))
```

Predict using the decision tree model on the test data

```
predicted_values2 <- predict(tree2, newdata = TestingSet2, type = "class")
```

Assuming 'is_fraud' is the actual target variable in the test data

```
actual_values2 <- TestingSet2$is_fraud
actual_values2 <- as.factor(actual_values)
levels(actual_values2) <- c("0", "1")
```

```
tree_predictions2 <- predict(tree2, TestingSet2, type = 'class')
head(tree_predictions2)
```

```
##  2  8 13 14 27 30
##  0  0  0  0  0  0
## Levels: 0 1
```

Create confusion matrix

```
conf_matrix2 <- confusionMatrix(predicted_values2, actual_values2)
```

```
print(conf_matrix2)
```

```
## Confusion Matrix and Statistics
```

```
##
##           Reference
## Prediction      0      1
##           0 138338    154
##           1    58    379
##
```

```
##           Accuracy : 0.9985
##           95% CI : (0.9983, 0.9987)
##      No Information Rate : 0.9962
##      P-Value [Acc > NIR] : < 2.2e-16
##
```

```
##           Kappa : 0.7807
##
```

```
##      McNemar's Test P-Value : 6.817e-11
##
```

```
##           Sensitivity : 0.9996
##           Specificity : 0.7111
##           Pos Pred Value : 0.9989
```

```
##          Neg Pred Value : 0.8673
##          Prevalence : 0.9962
##          Detection Rate : 0.9957
## Detection Prevalence : 0.9969
##          Balanced Accuracy : 0.8553
##
##          'Positive' Class : 0
##
```

5.3 KNN

5.3.1 KNN Balanced

```
fraudtrain1 <- TrainingSet1$is_fraud

#k selection with sqrt of total obs
print(sqrt(nrow(fraudtrain1)/2))

## numeric(0)

k <- 455

knnmodel1 <- knn(train = TrainingSet1, test = TestingSet2, cl = fraudtrain1,
k = k)

accuracy <- 100 * sum(TestingSet2$is_fraud == knnmodel1) / nrow(TestingSet2)
cat("k =", k, "Accuracy =", accuracy, "\n")

## k = 455 Accuracy = 89.69258

# Get predicted labels from the KNN model
predicted_labels1 <- as.numeric(knnmodel1)

# Calculate confusion matrix
conf_matrix1 <- table(Actual = TestingSet2$is_fraud, Predicted =
predicted_labels1)

# Calculate True Positives (TP), False Positives (FP), True Negatives (TN),
False Negatives (FN)
TP <- conf_matrix1[2, 2]
FP <- conf_matrix1[1, 2]
TN <- conf_matrix1[1, 1]
FN <- conf_matrix1[2, 1]

# Calculate True Positive Rate (TPR) and False Positive Rate (FPR)
TPR <- TP / (TP + FN)
TNR <- TN / (TN+FP)
FPR <- FP / (FP + TN)

TPR
```



```
## [1] 0.9043152
```

```
TNR
```

```
## [1] 0.8968973
```

```
FPR
```

```
## [1] 0.1031027
```

5.3.2 KNN Unbalanced Dataset

```
fraudtrain2 <- TrainingSet2$is_fraud
```

```
#k selection with sqrt of total obs
```

```
print(sqrt(nrow(fraudtrain2)/2))
```

```
## numeric(0)
```

```
k <- 455
```

```
knnmodel2 <- knn(train = TrainingSet2, test = TestingSet2, cl = fraudtrain2,  
k = k)
```

```
accuracy <- 100 * sum(TestingSet2$is_fraud == knnmodel2) / nrow(TestingSet2)
```

```
cat("k =", k, "Accuracy =", accuracy, "\n")
```

```
## k = 455 Accuracy = 99.61635
```

```
# Get predicted labels from the KNN model
```

```
predicted_labels2 <- as.numeric(knnmodel2)
```

```
# Calculate confusion matrix
```

```
conf_matrix2 <- table(Actual = TestingSet2$is_fraud, Predicted =  
predicted_labels2)
```

5.4 Random Forest

5.4.1 RF Balanced Dataset

```
#First step in running rf is converting target variable to factor
```

```
TrainingSet1$is_fraud <- as.factor(TrainingSet1$is_fraud)
```

```
# converting TestingSet$popular to factor
```

```
is_fraud_factor1 <- as.factor(TestingSet2$is_fraud)
```

```
# Assuming your data frame is called 'df' and the target variable is 'target'
```

```
rf_model1 <- randomForest(is_fraud ~ ., data = TrainingSet1, ntree = 100)
```

```
rf_model1
```

```
##
```

```
## Call:
```

```
## randomForest(formula = is_fraud ~ ., data = TrainingSet1, ntree = 100)
```

```
## Type of random forest: classification
```

```

##               Number of trees: 100
## No. of variables tried at each split: 5
##
##           OOB estimate of  error rate: 0.02%
## Confusion matrix:
##           0      1  class.error
## 0 249791      77 0.0003081627
## 1      0 166922 0.0000000000

rf_predictions1 <- predict(rf_model1, TestingSet2)
head(rf_predictions1)

##  2  8 13 14 27 30
##  0  0  0  0  0  0
## Levels: 0 1

cf_rf1 <- confusionMatrix(rf_predictions1, is_fraud_factor1)
cf_rf1

## Confusion Matrix and Statistics
##
##              Reference
## Prediction      0      1
##              0 138360      0
##              1      36    533
##
##              Accuracy : 0.9997
##              95% CI : (0.9996, 0.9998)
##              No Information Rate : 0.9962
##              P-Value [Acc > NIR] : < 2.2e-16
##
##              Kappa : 0.9672
##
##  Mcnemar's Test P-Value : 5.433e-09
##
##              Sensitivity : 0.9997
##              Specificity : 1.0000
##              Pos Pred Value : 1.0000
##              Neg Pred Value : 0.9367
##              Prevalence : 0.9962
##              Detection Rate : 0.9959
##              Detection Prevalence : 0.9959
##              Balanced Accuracy : 0.9999
##
##              'Positive' Class : 0
##

```

5.4.2 RF Unbalanced Dataset

#First step in running rf is converting target variable to factor
 TrainingSet2\$is_fraud <- as.factor(TrainingSet2\$is_fraud)

```

# converting TestingSet$popular to factor
is_fraud_factor2 <- as.factor(TestingSet2$is_fraud)

# Assuming your data frame is called 'df' and the target variable is 'target'
rf_model2 <- randomForest(is_fraud~ ., data = TrainingSet2, ntree = 100)
rf_model2

##
## Call:
## randomForest(formula = is_fraud ~ ., data = TrainingSet2, ntree = 100)
##           Type of random forest: classification
##           Number of trees: 100
## No. of variables tried at each split: 5
##
##           OOB estimate of  error rate: 0.15%
## Confusion matrix:
##           0    1 class.error
## 0 415155    23 5.539793e-05
## 1    584 1028 3.622829e-01

rf_predictions2 <- predict(rf_model2, TestingSet2)
head(rf_predictions2)

##  2  8 13 14 27 30
##  0  0  0  0  0  0
## Levels: 0 1

cf_rf2 <- confusionMatrix(rf_predictions2, is_fraud_factor2)
cf_rf2

## Confusion Matrix and Statistics
##
##              Reference
## Prediction      0      1
##           0 138390    193
##           1      6    340
##
##              Accuracy : 0.9986
##              95% CI : (0.9984, 0.9988)
##      No Information Rate : 0.9962
##      P-Value [Acc > NIR] : < 2.2e-16
##
##              Kappa : 0.7729
##
##  Mcnemar's Test P-Value : < 2.2e-16
##
##              Sensitivity : 1.0000
##              Specificity : 0.6379
##      Pos Pred Value : 0.9986
##      Neg Pred Value : 0.9827
##              Prevalence : 0.9962

```

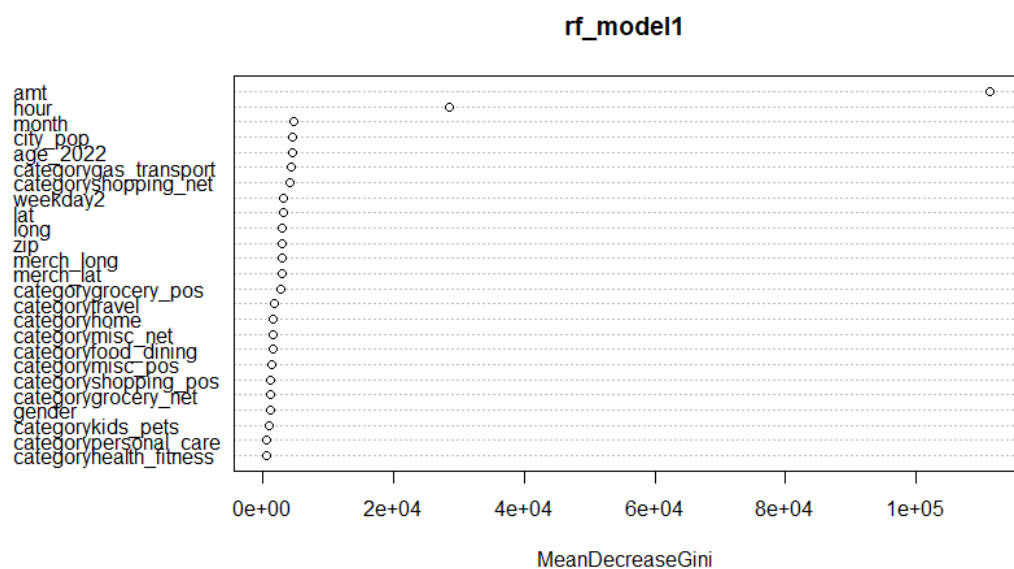
```
##          Detection Rate : 0.9961
##    Detection Prevalence : 0.9975
##          Balanced Accuracy : 0.8189
##
##          'Positive' Class : 0
##
```

Calculating Variable Importance

```
importance_values <- importance(rf_model1)
importance_values
```

```
##                               MeanDecreaseGini
## categoryfood_dining           1503.8373
## categorygas_transport         4351.5447
## categorygrocery_net          1207.6752
## categorygrocery_pos          2821.2147
## categoryhealth_fitness        509.4334
## categoryhome                  1605.8688
## categorykids_pets              896.9427
## categorymisc_net              1530.7656
## categorymisc_pos              1360.5457
## categorypersonal_care          631.2534
## categoryshopping_net          4091.5001
## categoryshopping_pos          1232.9741
## categorytravel                1668.5266
## amt                           111257.4079
## gender                        1205.6115
## zip                           2990.3055
## lat                           3160.2742
## long                           3018.1267
## city_pop                       4486.7403
## merch_lat                      2884.2972
## merch_long                     2889.8851
## age_2022                       4468.6060
## hour                          28536.1991
## weekday2                       3192.7751
## month                          4737.7620
```

```
varImpPlot(rf_model1)
```



6 Model Evaluation

Models are evaluated using accuracy from confusion matrix, testing error and also AUC score.

6.1 Plotting ROC Curves

#converting prediction scores data type before plotting curves

```
cls_prediction2 <- as.numeric(cls_prediction2)
tree_curve2 <- as.numeric(tree_predictions2)
rf_predictions2 <- as.numeric(rf_predictions2)
```

#calculating AUC curves of models# Calculate ROC and AUC using pROC

```
cls_score2 <- print(paste('cls roc_roc_curve score is',auc(cls_roc_curve2)))
```

```
## [1] "cls roc_roc_curve score is 0.499877154315858"
```

```
tree_score2 <- print(paste('tree_roc_curve score is',auc(tree_roc_curve2)))
```

```
## [1] "tree_roc_curve score is 0.995440257253315"
```

```
rf_score2 <- print(paste('rf_roc_curve score is', auc(rf_roc_curve2)))
```

```
## [1] "rf_roc_curve score is 0.805114540848315"
```

```
#creating the ROC function
```

```
cls_roc_curve2 <- roc(TestingSet2$is_fraud, cls_prediction2)
```

```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls < cases
```

```
tree_roc_curve2 <- roc(TestingSet2$is_fraud, tree_curve2)
```

```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls < cases
```

```
rf_roc_curve2 <- roc(TestingSet2$is_fraud, rf_predictions2)
```

```
## Setting levels: control = 0, case = 1
```


```
## Setting direction: controls < cases
```

```
# Plotting ROC Curves
```

```
plot(cls_roc_curve2, col = "blue", print.auc = TRUE)
```

```
plot(tree_roc_curve2, col = "red", add = TRUE)
```

```
plot(rf_roc_curve2, col = "yellow", add = TRUE)
```



6.2 Table of Results

```
# Create a new table with some sample data
Model_Comparison <- data.frame(
  Model = c("Logistic Regression", "Logistic Regression", "Decison Trees",
"Decision Trees", "Random Forests", "Random Forests", "KNN", "KNN"),
  Dataset = c("Balanced", "Unbalanced", "Balanced",
"Unbalanced", "Balanced", "Unbalanced", "Balanced", "Unbalanced"),
  Accuracy = c(0.861, 0.996, 0.994, 0.999, 0.999, 0.999, 0.897, 0.996),
  TestingError = c(0.139, 0.004, 0.006, 0.001, 0.103, 0.004, 0.001, 0.001),
  Sensitivity = c(0.999, 0.996, 0.991, 0.999, 0.904, 0.722, 0.999, 0.999),
  Specificity = c(0.05, 0.00, 1, 0.711, 0.897, 0.826, 1, 0.647))
```

```
# Display the new table
print(Model_Comparison)
```

##	Model	Dataset	Accuracy	TestingError	Sensitivity	Specificity
## 1	Logistic Regression	Balanced	0.861	0.139	0.999	0.050
## 2	Logistic Regression	Unbalanced	0.996	0.004	0.996	0.000
## 3	Decison Trees	Balanced	0.994	0.006	0.991	1.000
## 4	Decision Trees	Unbalanced	0.999	0.001	0.999	0.711
## 5	Random Forests	Balanced	0.999	0.103	0.904	0.897
## 6	Random Forests	Unbalanced	0.999	0.004	0.722	0.826
## 7	KNN	Balanced	0.897	0.001	0.999	1.000
## 8	KNN	Unbalanced	0.996	0.001	0.999	0.647